

Multicore Programming

Prof. Dr. Gudula Rünger
Dr. Jens Lang

Professur Praktische Informatik
Technische Universität Chemnitz

Winter Semester 2014/2015

Organizational Information

- ▶ Time of **lecture**:
Monday 15:30 to 17:00, Room 1/305,
- ▶ Time of **tutorials**:
Wednesday 17:15 to 18:45, Room 1/375,
Wednesday 19:00 to 20:30, Room 1/375,
Tutor: Thomas Jakobs, M.Sc.
- ▶ **Exam**:
 - ▶ Date: 1st February, during the lecture hours.
 - ▶ There will be homework sheets for every tutorial.
 - ▶ **You have to solve 8 out of 12 homework sheets. A homework is considered as solved if 50 % of the answers are correct.**
- ▶ Requirements to obtain credit points:
 - Attendance and passing of the exam
- ▶ Website for the lecture:
http://www.tu-chemnitz.de/informatik/PI/lehre/plan/WS15/teach_mp.php.en


Overview

1. Definitions and Content of Lecture
2. A Short Overview: Multicore Processors
3. Parallel Programming Concepts
4. Thread Programming
5. OpenMP

Parallel Region
Parallel Loop
More Directives
Control parameters
Critical Section, Atomic Operations, Reductions

Multicore Processor

- ▶ Multiple execution cores are placed on one chip.
- ▶ Each core typically has a private level-1 cache.
- ▶ The cores of a multicore processor typically share a level-2 cache.
- ▶ The cores have access to a shared memory (off-chip).
- ▶ The cores use shared resources (cache, TLB, bus etc.)



Multicore Processor

- ▶ Delivery by Intel, AMD, SUN, IBM since 2005
- ▶ In 2015 typically: processors with 2 (dualcore) or 4 (quadcore) cores, sometimes, 6 or 8 cores
- ▶ Highly parallel multicore processors (MIC – many integrated cores):
 - ▶ Intel's Xeon Phi co-processor with 61 cores (based on 80586 Pentium architecture).
 - ▶ Phytium's Mars processor with 64 cores (based on ARMv8).

Software for Multicore Processors

- ▶ Programs should take advantage of the cores of a multicore processor, i.e. all cores should compute tasks in parallel
~> **Parallel Programming**
- ▶ The cores use a shared memory
~> **Shared-Memory Programming**
(in contrast to message-passing programming)
- ▶ The cores interact by using shared resources
~> **Threading Technology**
- ▶ Programming environments:
Pthreads, Java Threads, OpenMP, C++11

Contents of the Lecture

- ▶ Architecture of Multicore Processors
- ▶ Overview: Parallel Programming
- ▶ Concepts of Shared Memory Programming
- ▶ Programming with OpenMP
- ▶ Scheduling Methods for Threads
- ▶ Cache, Cache Consistency and Cache Coherency
- ▶ Methods for Optimizing Cache Accesses (Loop Fusion, Loop Tiling)
- ▶ GPU Programming with OpenCL and CUDA
- ▶ Pthread Programming
- ▶ Software Transactional Memory

Literature

1. Parallel Programming for Multicore and Cluster Systems Second Edition. Rauber, Rünger, Springer, 2013.
2. Parallele Programmierung 2. Auflage. Rauber, Rünger, Springer, 2013.
3. Multicore: Parallele Programmierung. Rauber, Rünger, Springer, 2008.
4. Intel Threading Building Blocks: Outfitting C++ for Multi-Core Processor Parallelism, Reinders, O'Reilly 2007.
5. OpenMP, Informatik im Fokus, Hoffmann, Lienhart, Springer, 2008.
6. Multi-Core Programming – Increasing Performance through Software Multi-Threading. Akhter, Roberts, Intel Press, 2006.
7. Using OpenMP – Portable Shared Memory Parallel Programming. Chapman, Jost, van der Pas, MIT Press, 2008.
8. OpenMP Application Program Interface, Version 4.0, 2013, OpenMP Architecture Review Board, <http://www.openmp.org>
9. CUDA programming – A developer's guide to parallel computing with GPUs, Shane, Cook, Waltham, Morgan Kaufmann, 2013.
10. Pthreads programming, Dick Buttlar et al., O'Reilly, 2013.

Detailed Literature

| Chapter | English Literature | German Literature |
|--|---|-------------------------------|
| Architecture: A short overview on multicore processors | [1] 2.1 -2.4 | [2] 2.1 - 2.4 [3] 1.1-1.4 |
| Parallel Programming Concepts | [1] 3.1, 3.2, 3.3.1-3.3.5 [3] 2.1 bis 2.4, 2.9 | [2] 3.1-3.5, 4.2.1-4.2.2.1 |
| Thread Programming | [1] 3.6-3.7 | [2] 6.1 [3] 3.1-3.5 |
| OpenMP Programming | [1] 6.3 | [2] 6.4 [3] 6 |
| Cache, Cache-consistency, -Coherency | [1] 2.7 | [2] 2.7 |
| Synchronization Methods | [1] 6.1,6.2.6 | [2] 6.3.1 [3] 5.1,5.6 |
| Software Transactional Memory | [1] 6.2.6 | [3] 7 |

Note: The numbers enclosed in brackets [] refer to the list of books on the literature slide.

Overview

1. Definitions and Content of Lecture

2. A Short Overview: Multicore Processors

3. Parallel Programming Concepts

4. Thread Programming

5. OpenMP

Parallel Region

Parallel Loop

More Directives

Control parameters

Critical Section, Atomic Operations, Reductions

OpenMP 3.5

Overview

2. A Short Overview: Multicore Processors

Evolution of Micro Processors

Parallelism on Processor Level

Architecture of Multicore Processors

Examples

Evolution of Micro Processors

- ▶ The development of modern computer architectures is strongly pushed by the rapid technological development of electrical circuits;
VLSI circuits (very large scale integrated circuits)
→ high complexity of modern computer systems.
- ▶ The reason for the rapid development is the need for **higher performance** of processors.
- ▶ Around 2005: transition from **megahertz era** to **multicore era**
→ processors do not become faster any more, but more parallel.
- ▶ Technological evolutions include parallelism on different levels.

Trends in the Development of Processor Chips

- ▶ A rough measure of the complexity and performance of a processor is the **number of transistors**
- ▶ **Moore's Law:** The **number of transistors** per chip **doubles** every **18 to 24 months**.
 - ▶ Observation by Gordon Moore (Co-Founder of Intel), now valid for about 40 years.
- ▶ **More detailed:**
The number of **transistors per chip area** increases by about 35 % per year. Additionally, the chip area increases between 10 % and 20 % per year.

Performance Evolution of Processors

- ▶ Increase of **processor performance** by 55% (integer operations) or 75% (floating point operations) per year (1980-98)
- ▶ Increase of **memory capacity** of DRAM memory chips by 40% to 60% per year since 1977.
- ▶ Decrease of **memory access latency** of DRAM memory chips on average by about 25% per year.
~> Problem of memory wall

Reasons for the previous performance increase

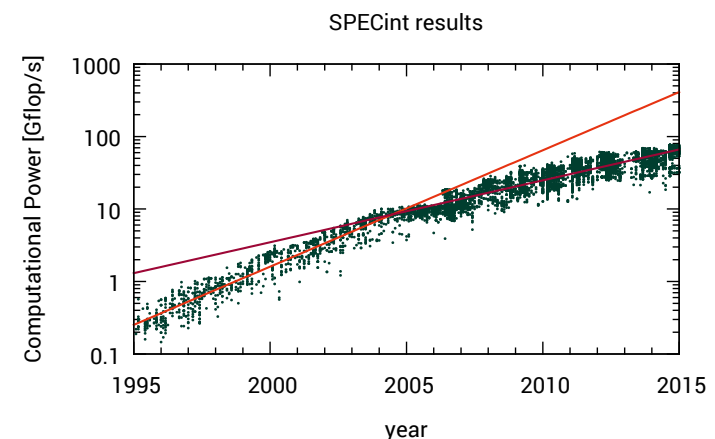
- (1) The **clock rate** of micro processors increases on average by about 30% per year
- (2) Architectural enhancements: Make use of the increasing number of transistors for **additional functionality** and **parallelism** on the instruction level (ILP - instruction level parallelism)

Architectonic Enhancements

Additional functionality and parallelism on the instruction level due to

- ▶ Increase of the **internal word length** to 32 or 64 bits;
- ▶ Make use of **internal pipelining** by decomposing the instruction processing into pipeline stages
~> Reduction of the **latency** per instruction (more instructions per time unit);
- ▶ Usage of **multiple functional units**, which receive instructions dynamically from a dispatch unit
~> Independent machine instructions can be executed in **parallel**
- ▶ Integration of **multi-level caches** onto the die of the processor
~> Reduction of the **average memory access latency**;
▶ reduces the Memory Wall Problem

Chart of Performance Increase



Each dot stands for a result of the SPECint benchmark with a CPU presented in the respective year.

Why a further performance increase is impossible

(1) Clock Rate:

- ▶ The **memory access latency** is not reduced to the same extent as the clock rate increases. This leads to a relative increase of the memory access time.
1990: Intel i486, 6–8 clock cycles per main memory access,
2006: Intel Pentium, more than 220 clock cycles per main memory access
- ▶ The increased transistor count is achieved by a higher packaging density, which results in an increased **heat generation** per chip area.
- ▶ An increased clock rate (time per clock cycle decreases) results in a **decreased distance** a signal can travel in one clock cycle. Therefore, several clock cycles are required for the complete processing of an instruction.
Layout avoiding long distances on the chip are more complex.

↪ A further increase of the clock rate is neither possible nor useful.

Why a further performance increase is impossible

(2) Instruction-level parallelism (pipelining; several functional units):

- ▶ Could be increased from a technological point of view;
- ▶ **Dependencies between instructions** prevent utilization of several functional units and deeper pipelines;
- ▶ The **control logic** requires a **high hardware complexity** and results in high energy requirements.

Note:

A lot of these (older) architectural enhancements require no or only minor **support from the compiler or operating system**

- ▶ Dynamic instruction scheduling managed by hardware;
- ▶ Cache management with reload strategy handled by hardware.

VLIW Processors (Very Long Instruction Word)

- ▶ **Basis:** The functional units are controlled by **static scheduling** by letting the **compiler** arrange independent instructions into long instruction words; ↪ the **parallel** processing of instructions is **explicitly** visible in the instruction words;
↪ EPIC (explicitly parallel instruction computers)
- ▶ **Examples:** Intel IA-64-Architecture (Itanium, Itanium 2);
Trimedia TM32-Architecture (processor for embedded systems);
Texas Instruments TMS320c6x (Digital Signal Processor);
- ▶ The **micro-processor hardware** is **simpler**, because the complexity for **dynamic scheduling** by the processor is omitted. The **compiler** can ensure by analysis methods that the functional units are efficiently utilized.
- ▶ **Requirement:** The compiler has to recognize independent instructions and has to arrange them into the instruction word

Overview

2. A Short Overview: Multicore Processors

Evolution of Micro Processors

Parallelism on Processor Level

Architecture of Multicore Processors

Examples

Parallelism on Processor Level

Simultaneous Multithreading or Hyperthreading Processors

- **Basis:** Hyperthreading is based on a **duplication of the processor region for the storage of the processor state** on the chip of the processors;
This includes: **general purpose registers, machine status registers, interrupt controller and its registers.**
- **Result:** The physical processor is seen as **two logical processors**
~> The operating system is able to assign two processes or threads to the processor; those may come from **the same** or **different applications**.
- The logical processors **share almost all resources** of the physical processor (caches, functional units, control logic, buses etc.)
~> Hyperthreading-Technology leads to an **increased chip area by only 5%.**

Hyperthreading Processors

- Advantage of hyperthreading: In case of **stalls** of the first logical processor, the **processing resources** can be used by the second logical processor.
~> **Increased utilization** of the processing resources.
- **Reasons for stalls:** cache misses, wrong branch prediction, dependencies between instructions, pipeline hazards etc.
- Depending on the application, Hyperthreading can **increase the runtime by 15% to 30%.**

Hyperthreading-Processors

- The operating system has to be able to control **two logical processors**. On systems with more than one processor, the operating system has to distinguish between **logical** and **physical** processors.
- More than one thread has to be available for execution;
~> the program has to be prepared accordingly;
- Case 1: **Multithreaded programming** (Pthreads, Java Threads)
- Case 2: **OpenMP** with compiler support (parallel loops, task queues, parallel sections)

Overview

1. A Short Overview: Multicore Processors
 - Evolution of Micro Processors
 - Parallelism on Processor Level
 - Architecture of Multicore Processors
 - Examples
2. A Short Overview: Multicore Processors


Multicore Processors

- ▶ Several **execution cores** are integrated into the chip of a processor.
Every execution core appears to the operating system as a separate **logical processor** with separate execution resources, which have to be controlled separately.
- ▶ **Design options** for Multicore-Processors:
 - ▶ Hierarchical design,
 - ▶ Pipeline design,
 - ▶ Network-based design.

Design Options for Multicore Processors (I)


Hierarchical Design

Several processor cores share several caches, with a **tree like** organization;
Examples: IBM Power6, Intel Core i7, AMD Opteron, Sun UltraSparc T1 (Niagara);



Design Options for Multicore Processors (II)

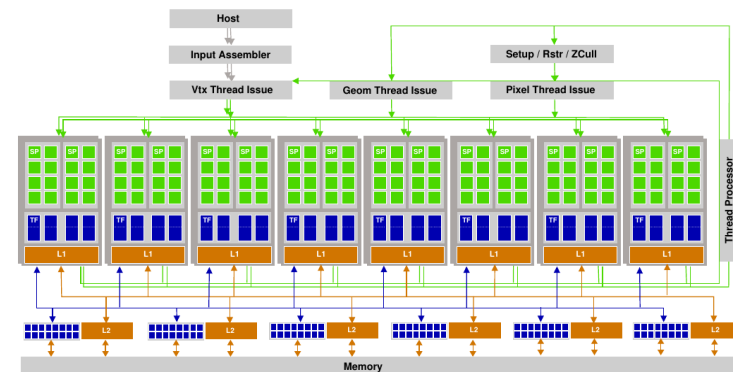
Hierarchical Design – Example I



Quad-core AMD Opteron (left) vs. quad-core Intel Xeon architecture (right) as examples for a hierarchical design.

Design Options for Multicore Processors (III)

Hierarchical Design – Example II




Nvidia GeForce 8800

Design Options for Multicore Processors (IV)

Pipeline Design

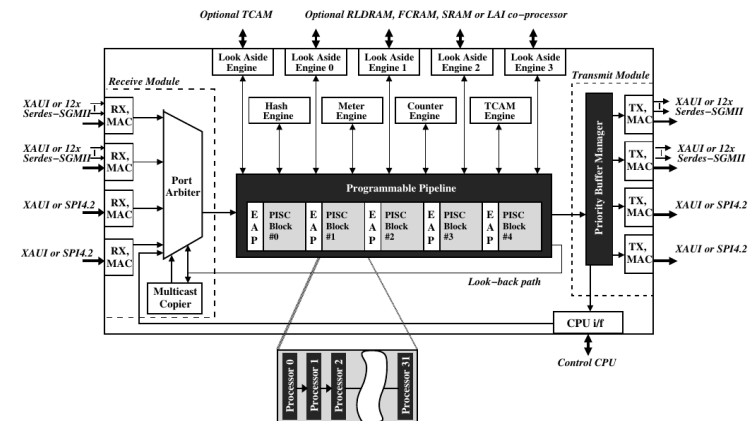
The data is processed **stepwise** core by core until the data reaches the last core, where the data is stored into memory.

Examples: GPUs, router processors for processing network packets; X10 and X11 processors from Xelerator: up to 200 separate VLIW cores;



pipelined design

Pipeline design - example




Xelerator X11 Network Processor

Design Options for Multicore Processors (V)

Network-based Design

The cores and their caches or memories are connected by an **interconnection network**;
→ all **data transfer** between processor cores is handled by the interconnection network.

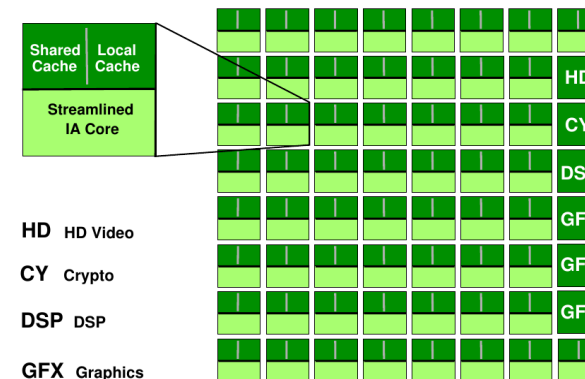
Example: Intel Teraflop chip with 80 processor cores organized in a 8×10 grid.



network-based design

Design Options for Multicore Processors (VI)

Network-based Design – Example



Intel Teraflop processor

Design Options for Multicore Processors (VII)

- ▶ **Evolution:** Intel Pentium D (without Hyperthreading) and Intel Pentium Extreme Edition (with Hyperthreading) introduced as dual-core in 2005;
AMD Dual Opteron also introduced in 2005;
2006: quad-core processors Intel Xeon 5300 (Clovertown);
Since 2010: oct-core processors.
- ▶ Multicore processors with **several heterogeneous processor cores** on a chip (e.g. ARM's big.LITTLE architecture);
- ▶ **Programming:** the processor cores of a chip can only be used by **threads executed in parallel**;
~~~> **Parallel programming techniques** have to be used.
- ▶ Difference to classical parallel systems: **Communication and synchronization** of the processor cores requires only a **small number of machine cycles** ~~~> **fine grained parallelism**;