


## Overview

### 2. A Short Overview: Multicore Processors

Evolution of Micro Processors  
Parallelism on Processor Level  
Architecture of Multicore Processors  
Examples


## Design Options for Multicore-Processors

Illustration of a **Intel Core i7** processor (hexa-core) from 2010



## Design Options for Multicore-Processors

Illustration of a **Pentium D** processor (dual-core) from 2005



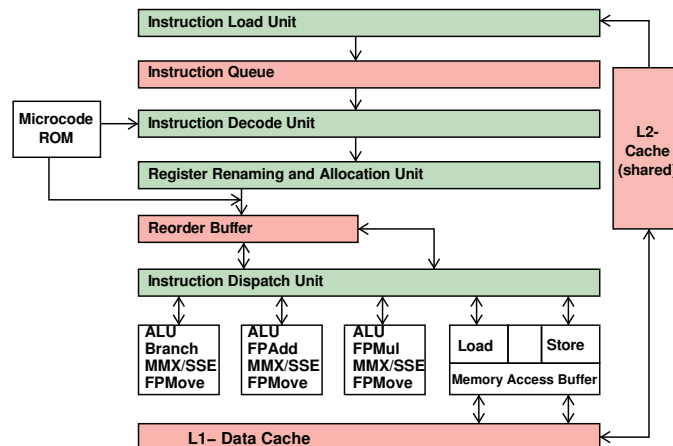
## Intel Core 2 Processors (I)

- ▶ Family of Intel processors with a common architecture
- ▶ Further development of the Pentium M (mobile) processor
- ▶ Features of the architecture:
  - ▶ Reduction of the pipeline length to max. 14 stages
  - ▶ Reduction of the clock rate  $\rightsquigarrow$  lower energy consumption
  - ▶ Power management unit for temporary shutdown of unused processor parts  
(reduction of energy consumption by reduction of clock rate or shutdown of L2 cache)
  - ▶ New SIMD extension (SSE)

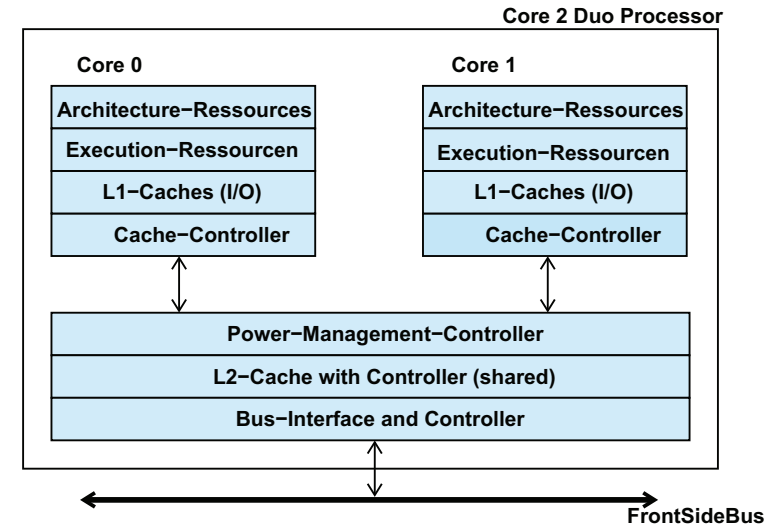
## Intel Core 2 Processors (II)

- ▶ Current versions: Core 2 Duo, Core 4 Duo (with 2 or 4 independent cores)
- ▶ No Hyperthreading (since Pentium M based)
- ▶ Separate L1 caches (32 KiB for instructions and 32 KiB for data)
- ▶ Shared L2 cache (4 MiB) for instructions and data
- ▶ Accessing the L2 cache (from cores or bus) is managed by the L2 controller.
- ▶ MESI protocol ensures coherency on the different levels of memory hierarchy (MESI – Modified, Exclusive, Shared, Invalid)
- ▶ Bus controller for data and I/O requests to/from the external bus
- ▶ Computations:
  - ▶ Individual instruction stream per core (incl. load/store)
  - ▶ Up to 4 instructions per core in parallel

## Instruction Processing and Memory Organization of a Core 2 Processor Core



## Core 2 Duo Architecture



## Overview

1. Definitions and Content of Lecture
2. A Short Overview: Multicore Processors
3. Parallel Programming Concepts
4. Thread Programming
5. OpenMP
6. Parallel Tasks
7. Pthread Programming

## Parallel Programming Concepts

- ▶ Parallel programming techniques are being used for many years, e.g. in
  - ▶ High-performance computing
  - ▶ computational science
 ↪ Multithreaded Programming is no new programming style
- ▶ Ubiquity of multicore processors in all areas of software development
- ▶ Essentials in multicore programming:
  - ▶ Transfer of parallel programming concepts
  - ▶ Providing of several instruction streams
- ▶ Subtopics of parallel programming:
  - Classification of parallel architectures, parallel programming models, parallel programming environments, parallel performance metrics

## Overview

### 3. Parallel Programming Concepts

Classification of Parallel Architectures

Design of Parallel Programs

Parallel Programming Models

Parallel Performance Measures

## Basic Concepts


- ▶ How to proceed when designing a parallel program?
- ▶ What properties of the parallel hardware should be made use of?
- ▶ What programming model should be used?
- ▶ How can the performance gain of the parallel program be quantified in comparison to the sequential program?
- ▶ Which parallel programming environment or language should be used?

## Flynn's Taxonomy

- ▶ Flynn's taxonomy characterizes parallel computers according to the **global control** and the resulting **data and control flows**.
- ▶ Four categories of parallel systems can be distinguished:
  - a) SISD – Single Instruction, Single Data: in each step **a single instruction** is applied to **a single data item**;
  - b) MISD – Multiple Instruction, Single Data: In each step **several different instructions** are applied to **the same data item**;
  - c) SIMD – Single Instruction, Multiple Data: In each step **the same instruction** is applied to **several data items**;  
(Many former parallel computers used that approach)
  - d) MIMD – Multiple Instruction, Multiple Data: In each step **different instructions** are applied to **different data items**;  
(Most of today's parallel computers)
 Further characterization by **memory organization**

## Memory Organization: Shared Address Space

- Multi-processor systems use an **interconnection network** for connecting the processors to a **global memory (shared memory machines)**:



- Data exchange** using **shared variables** and coordinated reading and writing → **explicit synchronization operations**
- Local and non-local memory accesses** are handled implicitly
- Variant: **symmetric multi-processors (SMP)**: The time for each **access to the shared memory** is the same for each processor and independent from the memory location (address).

## Interconnection Networks of Parallel Computers


- Used for **communication** in multi-computer systems
- Used for **memory access** in multi-processor systems

Features:

- Topology** of the interconnection network: **geometric structure** of the connections between the processors or memory modules;
  - Static** or **direct** interconnection networks: **The nodes are connected directly by using point-to-point wires.**
  - Dynamical** or **indirect** interconnection networks: **The processors and/or memory modules are indirectly connected using several wires and switches inbetween (bus based or switch based networks).**
- Routing techniques** – way of message transmission
  - Routing algorithm:** path selection
  - Switching strategy:** transmission mode

## Memory Organization: Distributed Address Space

- Multi-computer systems consist of **nodes (processors with local memory)**, connected to each other by an interconnection network:



- Exchange of data** is performed by **sending messages – message passing**;  
→ **Data exchange by explicit communication operations**
- Examples:** Blue Gene Project, IBM (gene research)  
PC clusters: Myrinet interconnect, Gigabit ethernet, Infiniband

## Overview

### 3. Parallel Programming Concepts

Classification of Parallel Architectures

Design of Parallel Programs

Parallel Programming Models

Parallel Performance Measures

## Creation of Parallel Instruction Streams (I)

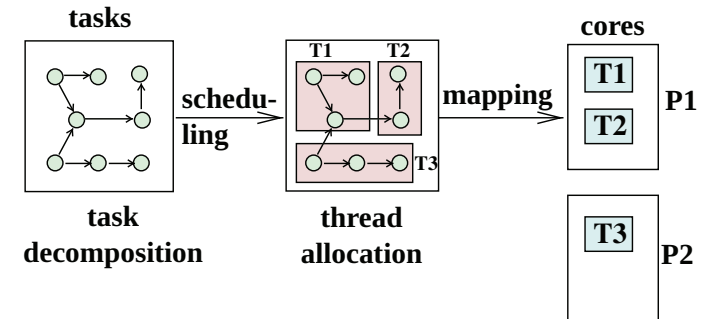
Basic idea: Creation of independent instruction streams for a parallel execution on several cores

- Division into **tasks**
  - Tasks are the basic units of parallelism
  - Task structure with task and control dependencies
  - The size of tasks is called **granularity**
- Mapping of tasks to threads or processes:  
**Scheduling** (static or dynamic)
- Mapping of processes or threads to cores  
(explicitly in the program or by the operating system)

## Execution of Parallel Tasks or Threads

- **Multitasking**: One processor, multitasking operating system, processes get time slices on the same processor, hiding of latencies (e.g. due to I/O)
- **Multiprocessing**: several physical processors, true parallel execution
- **Simultaneous Multithreading**: several logical processors on a physical processor, can decrease execution time of the tasks
- **Chip-Multiprocessing**: Fusion of multiprocessing and simultaneous multithreading:  
Multithreaded programs are not only executed concurrently, they are really executed in parallel on the processor

## Creation of Parallel Instruction Streams (II)



## Differences Between Multithreaded Programming for Multicores and Simultaneous Multithreading

- Performance improvements:  
SMT processors: Additional concurrent thread, e.g. for handling user requests  
Multicore: Whole application is parallelized.
- Impact of caches: **False Sharing** – two cores with separate caches work with data items that are different, but located in the same cache line.
- Thread Priorities:  
SMT: Thread with highest priority is executed first  
Multicore: Threads with different priorities can be executed in parallel  
↪ Can lead to varying results



## Overview

### 3. Parallel Programming Concepts

Classification of Parallel Architectures

Design of Parallel Programs

Parallel Programming Models

Parallel Performance Measures

## Parallel Programming Models (II)

- ▶ Should parallel program parts be specified explicitly?
  - ▶ Implicit parallelism: no specification of parallelism (e.g. parallelizing compilers, functional programming languages)
  - ▶ Explicit Parallelism
    - ▶ Implicit decomposition (Fortran Extensions)
    - ▶ Explicit decomposition (e.g. coordination languages like Linda)
    - ▶ Explicit communication and synchronization (e.g. MPI, Pthreads)
- ▶ How are parallel program parts specified?
  - ▶ **Thread:** Stream of instructions, which can be executed in parallel with other instruction streams (threads).
  - ▶ Processes which are executed on a separate physical processor or logical processor

## Parallel Programming Models (I)

**Abstract view** of the parallel system, which executes the parallel software, i.e.

programmers' view of the parallel hardware and system software like the operating system, parallel programming languages or libraries, compilers and runtime libraries

- ▶ On which level of a program should parallel program parts be executed?
  - ▶ Parallelism on instruction level
  - ▶ Parallelism on statement level
  - ▶ Parallelism on loop level, parallel loops
  - ▶ Parallelism on function level

## Parallel Programming Models (III)

- ▶ How is the parallel execution of program parts organized in relation to each other?
  - ▶ SIMD or SPMD (Single Program, Multiple Data)
  - ▶ Synchronous, asynchronous
  - ▶ Programming pattern: e.g. Pipelining, Master-Worker or Producer-Consumer-Model,
- ▶ How is the information exchange organized between parallel program parts?
  - ▶ Communication
  - ▶ Shared variables
- ▶ Which kind of synchronization can be used?
  - ▶ Barrier-Synchronization causes all threads to wait for each other
  - ▶ Coordination of the threads to avoid conflicts

## Overview

### 3. Parallel Programming Concepts

Classification of Parallel Architectures

Design of Parallel Programs

Parallel Programming Models

Parallel Performance Measures

## Costs of Parallel Programs

- ▶ The **costs**  $C_p(n)$  of a parallel program  $P$  with input size  $n$  on  $p$  processors is the **total time**, that the involved processors require for the execution of  $P$ :

$$C_p(n) = T_p(n) \cdot p$$

- ▶ The cost of a parallel program is a **measure for all the computations performed**.
- ▶ A parallel program is **cost optimal** if  $C_p(n) = T^*(n)$  holds, where  $T^*(n)$  is the runtime of the fastest sequential program;  
A cost optimal program requires as many computations as the fastest sequential program.  
**Difficulty:** The fastest sequential program or method is possibly not known or can only be determined with high efforts.
- ▶ The costs is often called **work** or **Processor-Time-Product**;

## Parallel Runtime of a Program

- ▶ The **parallel runtime**  $T_p(n)$  of a parallel program  $P$  with input size  $n$  on  $p$  processors is the **time between the start of program**  $P$  and the **termination** of the computations of  $P$  on **all** processors.
- ▶ For computers with physically **distributed memory**  $T_p(n)$  consists of:
  - ▶ Time for **local computations**
  - ▶ Time for **data exchange** with communication operations;
  - ▶ **Waiting time** of processors e.g. because of load imbalance
  - ▶ Time for **synchronization** of the executing processors or a subset of the executing processors
- ▶ For computers with shared memory the time for **data exchange** is replaced by the time for the **access to global data**.

## Speedup of a Parallel Program

- ▶ The **Speedup**  $S_p(n)$  of a parallel program  $P$  with input size  $n$  on  $p$  processors is defined as:

$$S_p(n) = \frac{T^*(n)}{T_p(n)} \quad \frac{\text{sequential time}}{\text{parallel time}}$$

- ▶  $T_p(n)$  = Parallel runtime of a parallel program  $P$  on  $p$  processors;
  - ▶  $T^*(n)$  = Runtime of an optimal sequential implementation for the solution of the problem;
- ▶ The speedup is a **measure of the relative speed increase compared to the best sequential implementation**
- ▶ Typically at most **linear speedup** can be reached:  
 $S_p(n) \leq p$  (theoretical upper bound)
- ▶ In practice, due to cache effects a **super linear speedup** can occur.

## Efficiency of a Parallel Program

- ▶ The **Efficiency**  $E_p(n)$  of a parallel program P with input size  $n$  on  $p$  processors is defined by

$$E_p(n) = \frac{T^*(n)}{C_p(n)} = \frac{S_p(n)}{p} = \frac{T^*(n)}{p \cdot T_p(n)}$$

- $C_p(n)$  = Parallel program cost
- $T_p(n)$  = Parallel runtime of a parallel program P
- $T^*(n)$  = Runtime of the best sequential implementation
- ▶ The efficiency is a **measure** for the portion of the runtime, that is required for computations that are also present in the sequential program.
- ▶ The ideal speedup  $S_p(n) = p$  is equivalent to  $E_p(n) = 1$ .

## Scalability

- ▶ The **scalability** of a parallel program on a parallel computer is a **measure** for the property **to achieve a performance increase proportional to the number  $p$  of processors**.
- ▶ Common observations:
  - For a fixed problem size  $n$  and an increasing number of processors  $p$  a **saturation of the speedup** occurs;
  - For a fixed number of processors  $p$  and an increasing problem size  $n$  an **increase of the speedup** occurs.
- ▶ **Concretization:** Scalability means, that the efficiency of a parallel program is constant when the number of processors  $p$  and the problems size  $n$  is increased

## Amdahl's Law

- ▶ When the parallel implementation requires a (constant) fraction  $f$ ,  $0 \leq f \leq 1$ , to be computed **sequentially**, the runtime of the parallel implementation is composed of:
  - The runtime  $f \cdot T^*(n)$  of the **sequential part** and
  - The runtime of the **parallel part**, which is at least  $(1 - f)/p \cdot T^*(n)$

- ▶ Attainable speedup

$$S_p(n) = \frac{T^*(n)}{f \cdot T^*(n) + \frac{1-f}{p} T^*(n)} = \frac{1}{f + \frac{1-f}{p}} \leq \frac{1}{f}.$$

- ▶ **Example:**  $f = 0.2 \rightarrow S_p(n) \leq 5$  independent from the number  $p$  of processors;

→ The sequential part has a big influence on the attainable speedup

for the efficient utilization of a high number of processors a