

Outline

1. Definitions and Content of Lecture
2. A Short Overview: Multicore Processors
3. Parallel Programming Concepts
4. Thread Programming
5. OpenMP
6. Parallel Tasks
7. Pthread Programming
8. GPU programming

OpenMP Standard

- ▶ **OpenMP** is a portable **standard** for the programming of **shared memory** systems.
- ▶ Standardization similar to the **MPI** standard for the distributed address space programming originally designed in 1997
- ▶ Naming: Open – *open standard*; MP – *multiprocessing*
- ▶ OpenMP provides an API (application programming interface) for C, C++ and Fortran
- ▶ **OpenMP Application Program Interface**, Version 2.5, May 2005; Version 3.0, May 2008; Version 4.0, July 2013
- ▶ OpenMP extends sequential programming languages by **language constructs**, **library functions** and **environment variables**
- ▶ More information: <http://www.openmp.org>

Literature

- Parallel Programming for Multicore and Cluster Systems, Rauber, Rünger, Springer 2010
- Parallele Programmierung, Rauber, Rünger, Springer 2008
- Multicore: Parallele Programmierung, Rauber, Rünger, 2008
- OpenMP, Informatik im Fokus, Hoffmann, Lienhart, Springer, 2008.
- Using OpenMP - Portable Shared Memory Parallel Programming. Barbara Chapman, Gabriele Jost, Ruud van der Pas, MIT Press, 2008.
- OpenMP Application Program Interface, Version 2.5, May 2005
- OpenMP Application Program Interface, Version 3.0, May 2008, OpenMP Architecture Review Board, <http://www.openmp.org>

SPEC OMP (OpenMP Benchmark Suite):

<http://www.spec.org/hpg/omp>

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Coordination and Synchronization

Programming Model

- ▶ An OpenMP program consists of a set of **cooperating threads**, created and destroyed by a **fork-join pattern**.
 - ▶ The execution of an OpenMP program begins with a **master thread**.
 - ▶ The master thread works sequentially until a `parallel` construct is encountered. At the parallel construct, the master thread creates a **team of threads** and becomes the **master** of that team.
- ▶ All threads execute the **same program text** (SPMD, specified in the `parallel` construct).
- ▶ At the end of the `parallel` **construct**, all threads perform an **implicit barrier synchronization**.
- ▶ All threads can access the **shared address space**
 ~> **synchronization** on **concurrent accesses** is needed

OpenMP Programming

- ▶ For using OpenMP, the **header file** `omp.h` is needed.
- ▶ There are constructs for the work sharing between parallel **threads**.
- ▶ There are constructs for the synchronization of parallel threads.
- ▶ Declaration of **shared** and **private variables**
Important: The **programmer** has to provide a **correct parallel program specification**, the compiler does not verify the correctness.

A First Example

A vector array of length 1000 is initialized **in parallel** by multiple threads:

```
1  const int n = 1000;
2  int array[size];
3  #pragma omp parallel for
4  for (int i = 1 ; i < n ; ++i)
5      array[i] = i;
```

Compiler directive (also called **pragma**):

```
#pragma omp parallel for
```

~> the statement (Line 4) after the `for` loop is executed in parallel, i.e. the statement `array[i] = i;` is executed in parallel.

- ▶ Implicit assignment of loop indices to threads

General Form of Compiler Directives

```
#pragma omp <directive> [clauses [[,] clauses] ...]
```

- ▶ A compiler directive has to end with a line break.
- ▶ The specification of clauses is optional.
- ▶ Clauses influence the behavior of directives.
- ▶ Different clauses are used for different directives.

OpenMP Compiler Support

- ▶ GCC supports OpenMP since Version 4.2
command line compiler switch **-fopenmp**.
- ▶ Intel C++ Compiler
supports the OpenMP standard and additionally Intel specific
directives (since Version 8).
- ▶ Sun Studio for Solaris OS supports OpenMP since Version 2.5.

The compiler variable `_OPENMP` tells, whether the compiler supports
OpenMP and OpenMP support is switched on.

Error Handling in OpenMP

- ▶ When an error occurs within an OpenMP directive, the
execution of the program is stopped.
- ▶ Program behavior is undefined according to the OpenMP
standard description in situations like:
for a parallel region no new threads are started, or a critical
section is not protected by synchronization.
- ▶ An explicit error message is not necessarily provided.

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Parallel Region

Parallel Loop

More Directives

Control parameters

Coordination and Synchronization

Work Sharing Constructs

- ▶ Parallelization of `for` loops
- ▶ Parallelization of non-iterative parallel regions (`section`)
- ▶ Parallel `task` (new concept in OpenMP 3.0)
- ▶ Execution explicitly not in parallel (`single`)

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Parallel Region

Parallel Loop

More Directives

Control parameters

Coordination and Synchronization

Parallel Region

For the **definition of a parallel region** the following directive is used:

```
#pragma omp parallel [parameter [parameter] ...]
    structured block of statements;
```

- ▶ This directive creates a **team of threads**
- ▶ The **master** of the team is the thread that calls the directive
- ▶ Master thread: thread number 0, remaining threads: consecutive thread numbers starting at one
- ▶ Each thread of the team (including the master) executes the **code block after the directive**.
- ▶ After the execution of the code block all threads except for the master are **destroyed**.
- ▶ The **number of threads in a team** is specified with the clause: `num_threads (expression)`
The expression has to return a positive integer value.

Parallel Region: Shared/Private Variables

The following optional **parameters** can be used for the declaration of **shared** or **private** variables:

- * **Private variables** are declared through:

```
private (list_of_variables)
```

Effect: Creation of **uninitialized copies** of the listed variables on the **runtime stack** of each thread.

Every thread accesses only its own **local copy**.

- * **Shared Variables** are declared through:

```
shared (list_of_variables)
```

Effect: Every thread of the team will access the **same variable** (same memory position) in the parallel region.

Parallel Region: Default Parameter

The **default parameter** is used for assigning a **default value** (shared or private) in a parallel region:

- ▶ `default (shared)`
specifies that all variables are **per default shared variables**.
- ▶ `default (none)`
specifies that each variable has to be **explicitly declared shared or private**.

More OpenMP Functions

- Determination of the **number of threads in a team**:
Runtime function:
`int omp_get_num_threads()`
- Obtain the unique **thread identifier** of the calling thread using a **consecutive numbering** of threads:
`int omp_get_thread_num()`
The **master thread** has the number **0**.
- **Nesting parallel regions** is possible.
Default: Only **one thread** executes the **parallel region at the inner-most nesting level**.
The call of
`void omp_set_nested (int nested)`
with `nested != 0` allows for the usage of **more than one thread for the parallel region on the inner-most nesting level**.

Example: Parallel Array Calculations

Array x is distributed to the threads of a parallel region:

```
#include <stdio.h>
#include <omp.h>
int npoints, iam, np, mypoints;
double *x;
int main() {
    scanf("%d", &npoints);
    x = (double *) malloc(npoints * sizeof(double));
    initialize();
    #pragma omp parallel shared(x,npoints) private(iam,np,mypoints)
    {
        np = omp_get_num_threads();
        iam = omp_get_thread_num();
        mypoints = npoints / np;
        compute_subdomain(x, iam, mypoints);
    }
}
```

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Parallel Region

Parallel Loop

More Directives

Control parameters

Coordination and Synchronization

Parallel Loop (I)

Inside a parallel region, iterations of **parallel loops** can be distributed to threads of the team of the parallel region:

```
#pragma omp for [parameter [parameter] ...]
for (i = lower_bound; i op upper_bound; incr_expr) {
    //loop body
}
```

- The **iterations** of the loop have to be **independent of each other**.
- The **number of iterations** must be known in advance.
- It is **forbidden to modify the iteration variable i** in the **loop body**.
In the loop body, i is treated as a **private variable** of the threads involved.

Parallel Loop (II)

- ▶ The expression `op` denotes a **boolean operator** from the set $\{<, <=, >, >=\}$.
- ▶ The expression `incr_expr` can be one of the following: `++i, i++, --i, i--, i=i+incr, i=incr+i, i=i-incr`
- ▶ At the end of a parallel loop, all participating threads are **synchronized implicitly**, which can be **avoided** by using the parameter `nowait`.
- ▶ The nesting of **for directives** within a parallel region is **forbidden**. But it is possible to nest **parallel regions**.

Scheduling of Loop Iterations (I)

The **mapping of loop iterations** to threads of a parallel region can be specified precisely by using the **scheduling parameter**:

- * `schedule (static, block_size):`
Static distribution – the iterations are assigned in **blocks** of size `block_size` in a **round-robin** fashion to the threads of the team. If `block_size` is not given, blocks of **almost equal size** are formed and assigned in a **blockwise distribution**.
- * `schedule (dynamic, block_size):`
Dynamic distribution of **blocks** of size `block_size` to threads. A thread gets a **new block** as soon as the thread finishes the computation of the previously assigned block. When the block size is not specified, **blocks of the size of 1** are used.

Scheduling of Loop Iterations (II)

- * `schedule (guided, block_size):`
Dynamic scheduling of blocks with **decreasing size**. Iterations are assigned in blocks of size

$$\text{block size} = \frac{\text{number of remaining iterations}}{\text{number of threads}}$$

For `block_size > 1`, a block contains **never less** than `block_size` iterations, except for the last block.

Default: `block_size = 1`;

- * `schedule (runtime):`
The **scheduling is specified at runtime**
Control: **environmental variable** `OMP_SCHEDULE`

Example: Matrix Multiplication with For Directive (I)

Static scheduling: every thread computes a **row block** of the matrix
Realization in **two steps** (initialization and calculation) with **implicit synchronization**.

```
#include <omp.h>
```

```
double MA[100][100], MB[100][100], MC[100][100];
int i, row, col, size = 100;
```

Example: Matrix Multiplication with For Directive (II)

```
int main() {
    read_input(MA, MB);
    #pragma omp parallel shared(MA,MB,MC,size) private(row,col,i) {
        #pragma omp for schedule(static)
        for (row = 0; row < size; row++) {
            for (col = 0; col < size; col++)
                MC[row][col] = 0.0;
        }
        #pragma omp for schedule(static)
        for (row = 0; row < size; row++) {
            for (col = 0; col < size; col++)
                for (i = 0; i < size; i++)
                    MC[row][col] += MA[row][i] * MB[i][col];
        }
    }
    write_output(MC);
}
```

Example: Matrix Multiplication – Double Parallel Loop

```
int main() {
    read_input(MA, MB);
    #pragma omp parallel private(row,col,i) {
        #pragma omp for schedule(static)
        for (row = 0; row < size; row++) {
            #pragma omp parallel shared(MA, MB, MC, size) {
                #pragma omp for schedule(static)
                for (col = 0; col < size; col++) {
                    MC[row][col] = 0.0;
                    for (i = 0; i < size; i++)
                        MC[row][col] += MA[row][i] * MB[i][col];
                }
            }
        }
    }
    write_output(MC);
}
```

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Parallel Region

Parallel Loop

More Directives

Control parameters

Coordination and Synchronization

Non-Iterative Worksharing Constructs

Computations of the parallel section can be **assigned explicitly** to threads, with the **section directive**:

```
#pragma omp sections [parameter [parameter]...]
{
    [#pragma omp section]
        structured block
    [#pragma omp section]
        structured block
    ...
}
```

The **sections** directive ends with an **implicit synchronization**.

Single Execution

If instructions in a parallel region may **only** be **executed once**, the **single** can be used:

```
#pragma omp single [parameter[parameter]...]
    structured block
```

Effect: The structured block is **executed by only one thread** of the team.

Syntactic Abbreviations

A parallel region with **only one for directive** can be abbreviated with:

```
#pragma omp parallel for [parameter[parameter]...]
    for (i = lower_bound; i op upper_bound; incr_expr) {
        loop body
    }
```

A parallel region with **only one sections directive** can be abbreviated with:

```
#pragma omp parallel sections [parameter[parameter]...]
{
    [#pragma omp section]
        structured block
    [#pragma omp section]
        structured block
    ...
}
```

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Parallel Region

Parallel Loop

More Directives

Control parameters

Coordination and Synchronization

Controlling the Number of Threads (I)

The **number of threads** used for a parallel region can be controlled with the help of **runtime functions**:

- Dynamic adjustment of the number of threads through the runtime system:

```
void omp_set_dynamic (int dynamic_threads)
```

- The value `dynamic_threads` $\neq 0$ uses a **dynamic adjustment**;
 - The value `dynamic_threads` $= 0$ **disables the dynamic adjustment**
- ↪ The runtime system will use the **specified number of threads**.

The function must be called **outside of all parallel regions**.

Controlling the Number of Thread (II)

- ▶ Query of the **current status**:
`int omp_get_dynamic (void)`
Return value 0, if no dynamic adjustment is set.
- ▶ **Set the number of threads** for the subsequent parallel regions:
`void omp_set_num_threads (int num_threads)`
- ▶ If dynamic thread adjustment is enabled: `num_threads` specifies the **maximum number of threads** to be used.
- ▶ If dynamic thread adjustment is disabled: `num_threads` denotes the **number of threads** in subsequent parallel regions.

Controlling the Nesting of Threads

- ▶ Controlling the **nesting in parallel regions**:
`void omp_set_nested (int nested)`
`nested = 0`: The inner parallel region is executed by one thread **sequentially**.
`nested ≠ 0`: The runtime system can use **more than one** thread for executing the inner parallel region.
- ▶ **Query the current status** of the nesting strategy:
`int omp_get_nested (void)`

Overview

5. OpenMP

General Information about OpenMP
Controlling Parallel Computations
Coordination and Synchronization
Critical Section, Atomic Operations, Reductions
Synchronization of Events
Synchronization using the OpenMP Runtime Library

Overview

5. OpenMP

General Information about OpenMP
Controlling Parallel Computations
Coordination and Synchronization
Critical Section, Atomic Operations, Reductions
Synchronization of Events
Synchronization using the OpenMP Runtime Library

Critical Section (I)

A **parallel region** is executed by multiple threads accessing the same **shared data**. Thus, there is the need for synchronization in order to protect critical sections for avoiding race conditions. OpenMP offers several possibilities:

- Definition of a **critical section**, which is executed by **only one thread** at a time:

```
#pragma omp critical [(name)]
structured block
```

Effect: When a thread encounters a critical section, it waits until no other thread executes the corresponding code section.

Critical Section (II)

- A critical section can be executed by only one thread at a time. This refers to all threads of the program, not only to the current team of threads.
- All code sections that are marked with **critical** are considered to be **one critical section**.
⇒ Exclusion of parallel execution of different code sections where synchronization is not necessary.
- Finer Granularity:
 - **Critical sections with names** (in round braces):
Threads have to wait for entering a **named critical section** until no other thread is executing the critical section with the same name.
 - **Unnamed critical sections** are handled like one critical section with an **unspecified** name.
- **Exiting** a structured block of a critical region with **break**, **continue** or **return** is **not allowed**. Otherwise, a proper release of a critical section is impossible.

Atomic Operations

An **atomic operation** is specified by:

```
#pragma omp atomic
statement
```

The statement must be of the following form: `x binop= E`
`x++, ++x, x--, --x` where `x` denotes a **variable** and `E` denotes a **scalar expression, not containing x**;
`binop` $\in \{+, -, *, /, \&, \wedge, \backslash, \ll, \gg\}$

Effect: After evaluating the expression `E`, the update of `x` becomes an **atomic operation**, i.e. a **non-interruptable operation**.

Note: The evaluation of `E` is **not an atomic operation**.

Global Reduction Operation

- **Global reduction operations** can be realized by the **reduction** clause:
`reduction (op:list)`
 with a reduction operator `op` $\in \{+, -, *, /, \&, \wedge, |, \&\&, ||\}$, `list` is a list of **reduction variables** (shared variables).
- **Effect:**
 - Each thread gets a **private copy** of each reduction variable. A thread can assign values to its private copy.
 - At the **end** of the surrounding region, a reduction operation of the local copies of each reduction variable is performed and stored into the corresponding **shared variable**.
 - This means, the shared result variable is computed by performing the given operation `op` with all local copies of the reduction variable.

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Coordination and Synchronization

Critical Section, Atomic Operations, Reductions

Synchronization of Events

Synchronization using the OpenMP Runtime Library

Barrier-Synchronization

An explicit **barrier synchronization** can be achieved by using the `pragma`:

```
#pragma omp barrier
```

Effect: All threads of the team wait at the barrier until all other threads reach it. Only when all threads have reached that barrier, the program continues.

Data Consistency

- ▶ A **consistent view** of the memory of shared data structures can be obtained explicitly by using

```
#pragma omp flush [(list)]
```

A **list of variables** (`list`) can be provided. Each of these variables gets a consistent value, i.e. all threads see the same value. If the list is empty, all variables get a consistent value.

- ▶ A **flush** directive is executed implicitly after executing:
 - ▶ a **barrier** directive,
 - ▶ when **entering** and after **leaving** a **critical** directive,
 - ▶ after **leaving** a **parallel** directive,
 - ▶ after leaving a **for**, **sections** or **single** directive, if no **nowait** parameter is given.

Example: Flush Directive

```
#pragma omp parallel private (iam, neighbor) shared (work, sync)
{
    iam = omp_get_thread_num();
    sync[iam] = 0;
    #pragma omp barrier
    work[iam] = do_work();
    #pragma omp flush (work)
    sync[iam] = 1;
    #pragma omp flush (sync)
    neighbor = (iam != 0) ? (iam - 1) : (omp_get_num_threads() - 1);
    while (sync[neighbor] == 0) {
        #pragma omp flush (sync)
    }
    combine (work[iam], work[neighbor]);
}
```

Master-Thread

Execution of a structured block within a parallel region **only by the master thread**:

```
#pragma omp master
    structured block
```

No other thread executes the specified structured block.

Overview

5. OpenMP

General Information about OpenMP

Controlling Parallel Computations

Coordination and Synchronization

Critical Section, Atomic Operations, Reductions

Synchronization of Events

Synchronization using the OpenMP Runtime Library

Simple Lock-Variables (I)

- ▶ Locking mechanism with **simple lock variables**
 - ▶ **simple lock variables** of type `omp_lock_t`
Simple lock variables can be **locked once**.
 - ▶ **Nestable lock variables** of type `omp_nest_lock_t`
Nestable lock variables can be **locked several times**.

- ▶ **Initialization** of lock variables:

```
void omp_init_lock (omp_lock_t *lock);
void omp_init_nest_lock (omp_nest_lock_t *lock);
```

Simple Lock Variables (II)

- ▶ **Destruction** of lock variables:

```
void omp_destroy_lock (omp_lock_t *lock)
void omp_destroy_nest_lock (omp_nest_lock_t *lock)
```

- ▶ **Setting** a lock variable:

```
void omp_set_lock (omp_lock_t *lock)
void omp_set_nest_lock (omp_nest_lock_t *lock)
```

Effect: The executing thread is **blocked until** the lock variable is **available**. Then, the lock variable is set to *locked* by this thread. A **simple lock variable** is available, when **no thread** currently holds it. A **nestable lock variable** is available, when at most the calling thread is holding it.

Simple Lock Variables (III)

- **Releasing** a lock variable:

```
void omp_unset_lock (omp_lock_t *lock)
void omp_unset_nest_lock (omp_nest_lock_t *lock)
```

A lock variable can only be **released** by the thread **holding** it.
A **nestable** lock variable **has to be released as many times** as it was **locked**.

- **Test** of a lock variable:

```
int omp_test_lock (omp_lock_t *lock)
int omp_test_nest_lock (omp_nest_lock_t *lock)
```

If the issued lock variable is available, it gets locked.
If the issued lock variable is locked, the calling thread is not blocked.

Return value: 1 on successful locking,
otherwise 0

Time measurement

- **Time measurement** in OpenMP:

```
double omp_get_wtime();
```

- Example:

```
{
    double d, time;
    ...
    d=omp_get_time();
    // part to measure
    time=omp_get_time()-d;
}
```